

Week 9: Elements of Compiler Design

Translation to Intermediate Code

Prof. Chris GauthierDickey
University of Denver

Generating LLVM IR

Generating LLVM IR

- ❖ We have made it to what is very near code generation
- ❖ From the AST, using Frame and our environment information, we will emit LLVM Intermediate Representation (IR) code
- ❖ This code can be compiled through the llvm compiler, which is freely available (and already lives on the Macs)

Why LLVM?

- ❖ We only have 10 weeks. We could go to a custom IR, and then optimize it, and then do instruction selection, but that's a ton of work in 2 weeks!
- ❖ However, the LLVM IR is now widely used, and from it you can generate code for almost every platform
- ❖ The LLVM IR was designed to be similar to an automatic grammar generator — if you generate the IR, you can compile and optimize for many platforms

The High Level

- ❖ LLVM divides code chunks into modules. A module is a translation unit (i.e., file)
- ❖ Each module has functions, global variables, and its own set of symbol table entries (note, LLVM handles the symbol tables automatically)
- ❖ Each module also has parts that help with binary code generation

Targets

```
; target data layout for Mac, change to m:w instead for Windows  
target datalayout = "e-m:o-i64:64-f80:128-n8:16:32:64-S128"
```

- ❖ The data layout specifies information about the architecture, in particular endianness, name mangling, floating point/integer sizes, and stack alignment
- ❖ The string consists of values separated by ‘-’
- ❖ e - little Endian, E - big Endian
- ❖ m:o - Mach-O name mangling, m:w - Windows COFF mangling
- ❖ i64:64 indicates the natural word size of the machine, 64 bits in this case, with a 64 bit alignment
- ❖ f:80:128 - floating point values of 80 and 128 bits. Note this means you can make floating point values of those sizes, 32-bit and 64-bit are always supported on all targets
- ❖ n:8:16:32:64 - valid integer sizes
- ❖ S128 - natural stack alignment, in bits (so in this case 16 bytes)

Targets

```
; target triple for Mac  
target triple = "x86_64-apple-macosx10.10.0"
```

- ❖ The triple tells it how to generate machine code—for which architecture in essence
- ❖ This triple is different for Mac and Windows (above is for the Mac)
- ❖ For testing on Windows, you can just leave this out, llvm will generate the default kind automatically, but warn you
- ❖ *Please do not include the target triple, or use the Mac target triple (if you're on a Mac) because I grade on a Mac!*

Types

- ❖ Throughout LLVM, you'll see that types are specified
 - ❖ *LLVM is statically, strongly typed!*
- ❖ This makes generating some kinds of code, like for a static link, a bit of a hassle
- ❖ But, the types allow for better optimizations in the end!

First Class Types

- ❖ Values of first class types *can only be produced by instructions*, these include:

Integer Types	iN (where N is the number of bits)	i8, i16, i32, i64, i1942652 (haha)	
Floating point types	half (16 bits) float (32 bits) double (64 bits)	fp128 (128 bits, 112-bit mantissa) x86_fp80 (80 bits)	ppc_fp128 (128-bit, 2 64-bit floating points)
Pointer Types	<type>*	i32*, i64*, float*	[4 x i32]* i32 (i32 *) *
Vector Types	< # elements > x <element type>	<4 x i32> <4 x float> <2 x i64>	
Label	Labels are just labels, but they're typed	mylabel:	

Other Types

Functions

<return type> (<parameter
list>)

i32 (i32)
float (i16, i32 *)
i32 (i8*, ...)

This last example is a varargs
call, such as printf

Arrays

[<# elements> x <type>]

[40 x i32]
[2 x [3 x [4 x i16]]]

Structures

%T1 = type { <type list> }
%T1 = type '<' { <type list> } '>'

%String = type { i64, i8* }
%triple = type <{i32, i32, i32}>

Void

No value and no size type

void

Basic Blocks

- ❖ LLVM code is divided into basic blocks
- ❖ Basic blocks begin with an (optional) Label and end with a terminating instruction (*ret, br, switch, indirectbr, invoke, resume, unreachable*)
- ❖ If you do not assign a label, one will be assigned for you, and these all look like incrementing numerical temporaries (e.g., `%1, %2, ...`)
- ❖ Basic blocks are important because we use them to jump around our code via their labels
- ❖ In addition, function bodies are a list of basic blocks

More on Basic Blocks

- ❖ Why even bother?
 - ❖ Basic blocks may be reordered for optimization
 - ❖ Functions simply contain a list of these without needed to generate a complicated graph with branching
 - ❖ You don't have to do anything special to define them, just use labels as necessary, and when you translate, you'll end up with terminator instructions

Identifiers

- ❖ LLVM have two basic types of identifiers, *global* and *local*.
 - ❖ Global identifiers are exactly how they sound: they're global variables accessible by all emitted code
 - ❖ Local identifiers are only accessible within the body of a function (so you can't use a local identifier from one function in a different function)

Identifiers

```
let
    ni x is 42
in
end
```

- ❖ Global identifiers begin with @

Identifiers

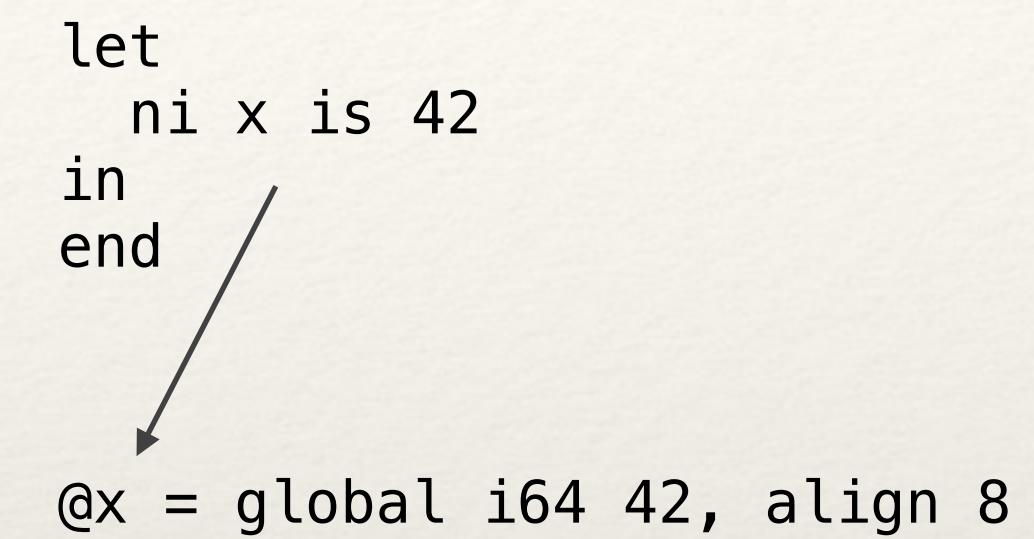
```
let
  ni x is 42
in
end
@x = global i64 42, align 8
```



- ❖ Global identifiers begin with @

Identifiers

```
let
  ni x is 42
in
end
@x = global i64 42, align 8
```



The diagram illustrates the scope of identifiers. A code snippet is shown with annotations:

- The identifier `x` is defined within the `let` block.
- The identifier `x` is used in the `ni` block.
- The identifier `x` is used in the `end` block.
- The identifier `@x` is defined outside the `let` block, serving as a global variable.

- ❖ Global identifiers begin with `@`
- ❖ Local identifiers begin with `%`

Identifiers

- ❖ Global identifiers begin with @
- ❖ Local identifiers begin with %

```
let
  ni x is 42
in
end
@x = global i64 42, align 8
```

```
let
  ni x is 42
in
let
  neewom addStuff(int y) as int is
    let
      ni z is 3
    in
      x + y + z
  in
    addStuff(6)
end
```

Identifiers

- ❖ Global identifiers begin with @
- ❖ Local identifiers begin with %

```
let
  ni x is 42
in
end
@x = global i64 42, align 8
```

```
let
  ni x is 42
in
let
  neewom addStuff(int y) as int is
    let
      ni z is 3
      in
        x + y + z
      in
        addStuff(6)
    end
  end
%z = alloca i64, align 8
store i64 3, i64* %z
```

Identifiers

- ❖ Notice we demonstrate two things:
we can load an integer constant into
a global easily

```
let
  ni x is 42
in
end
@x = global i64 42, align 8
```

```
let
  ni x is 42
in
let
  neewom addStuff(int y) as int is
    let
      ni z is 3
      in
        x + y + z
    in
    addStuff(6)
  end
%z = alloca i64, align 8
store i64 3, i64* %z
```

Identifiers

- ❖ Notice we demonstrate two things:
we can load an integer constant into
a global easily
- ❖ A local may be on the stack, and if it
is, it's an address with a pointer
type

```
let ni x is 42
in end
@x = global i64 42, align 8

let ni x is 42
in
let neewom addStuff(int y) as int is
    let
        ni z is 3
        in x + y + z
    in addStuff(6)
end

%z = alloca i64, align 8
store i64 3, i64* %z
```

Identifiers

- ❖ In addition, notice the ‘align 8’ specification
 - ❖ While integer types specify the number of bits, the align specification is in bytes
 - ❖ Thus, 64 bits is naturally aligned on 8 byte boundaries

```
let
  ni x is 42
in
end
@x = global i64 42, align 8
```

```
let
  ni x is 42
in
let
  neewom addStuff(int y) as int is
    let
      ni z is 3
      in
        x + y + z
      addStuff(6)
    end
  %z = alloca i64, align 8
  store i64 3, i64* %z
```

Stack vs Register

- ❖ Global variables are *always* pointers.
 - ❖ For example, @x global i64 42, align 8
 - ❖ This stores 42 in the global variable x
 - ❖ But what type is x? It's an i64 *, not an i64
 - ❖ You have to retrieve the value with a load instruction

Stack vs Register

- ❖ When we declare locals, we use `%` preceding the name
- ❖ However, we can declare two types of local variables, stack or frame resident
 - ❖ `%x = alloca i64, align 8`
 - ❖ This produces a value in the stack frame, we can store and load to and from it
 - ❖ `%x = add i64 42, 0`
 - ❖ This produces a *temporary* or register allocated variable, assigning 42 to it (because we just added two ints)

Stack vs Register

- ❖ If a variable escapes, you must place it on the stack
- ❖ If it doesn't, you may place it in a register
 - ❖ However, we have a major restriction, the LLVM IR uses Static Single Assignment (SSA) form
 - ❖ This means that every name can be used only once

Static Single Assignment

- ❖ Consider the following code
 - ❖ Obviously, the second declaration of x shadows the first (obvious to a human)

```
let
    ni x is 42
    ni x is 23
    ni y is 0
in
    now y is x
end
```

Static Single Assignment

- ❖ Now consider if we had to use unique names
- ❖ The compiler can easily optimize away the first x_1

```
let
  ni x is 42
  ni x is 23
  ni y is 0
in
  now y is x
end
```

```
let
  ni x1 is 42
  ni x2 is 23
  ni y1 is 0
in
  now y1 is x2
end
```

A more complicated example

- ❖ Ignoring the somewhat verbose language description, we have have to rename the variables so that the final value of y is the result of the branch
- ❖ Static single assignment only allows one assignment to y , so if we rename, we run into issues:

```
let
  ni x is 42
in
  now x is x - 7;
  let
    ni y is 3
    ni w is 0
  in
    if x <= 35 then
      now y is x - 3
    else
      now y is x - 10
    end;
    now w is y * 2;
    w
  end
end
```

A more complicated example

- ❖ First, let's rename x as x_1 , and x_2 (remember we're converting to LLVM IR, so while this program isn't valid in Ni, the equivalent IR would be fine)

```
let
  ni x1 is 42
in
  now x2 is x1 - 7;
  let
    ni y is 3
    ni w is 0
  in
    if x2 <= 35 then
      now y is x2 - 3
    else
      now y is x2 - 10
    end;
    now w is y * 2;
    w
  end
end
```

A more complicated example

- ❖ First, let's rename x as x_1 , and x_2 (remember we're converting to LLVM IR, so while this program isn't valid in Ni, the equivalent IR would be fine)
- ❖ Now let's rename w

```
let
  ni x1 is 42
in
  now x2 is x1 - 7;
  let
    ni y is 3
    ni w1 is 0
  in
    if x2 <= 35 then
      now y is x2 - 3
    else
      now y is x2 - 10
    end;
    now w2 is y * 2;
    w2
  end
end
```

A more complicated example

- ❖ First, let's rename x as x_1 , and x_2 (remember we're converting to LLVM IR, so while this program isn't valid in Ni, the equivalent IR would be fine)
- ❖ Now let's rename w
- ❖ Finally, rename y

```
let
  ni x1 is 42
in
  now x2 is x1 - 7;
  let
    ni y1 is 3
    ni w1 is 0
  in
    if x2 <= 35 then
      now y2 is x2 - 3
    else
      now y3 is x2 - 10
    end;
    now w2 is y? * 2;
    w2
  end
end
```

A more complicated example

- ❖ First, let's rename x as x_1 , and x_2 (remember we're converting to LLVM IR, so while this program isn't valid in Ni, the equivalent IR would be fine)
- ❖ Now let's rename w
- ❖ Finally, rename y
 - ❖ We're stuck! Should we put y_2 or y_3 on this line?

```
let
  ni x1 is 42
in
  now x2 is x1 - 7;
  let
    ni y1 is 3
    ni w1 is 0
  in
    if x2 <= 35 then
      now y2 is x2 - 3
    else
      now y3 is x2 - 10
    end;
    now w2 is y? * 2;
  w2
end
end
```

Introducing ϕ

- ❖ The ϕ instruction was designed to choose at that time which branch we just came from
- ❖ ϕ was so-named by IBM in the 80s when they developed the idea of SSA to stand for ‘phoney’, haha

```
let
  ni x1 is 42
in
  now x2 is x1 - 7;
  let
    ni y1 is 3
    ni w1 is 0
  in
    if x2 <= 35 then
      now y2 is x2 - 3
    else
      now y3 is x2 - 10
    end;
    now w2 is y? * 2;
    w2
  end
end
```

Introducing ϕ

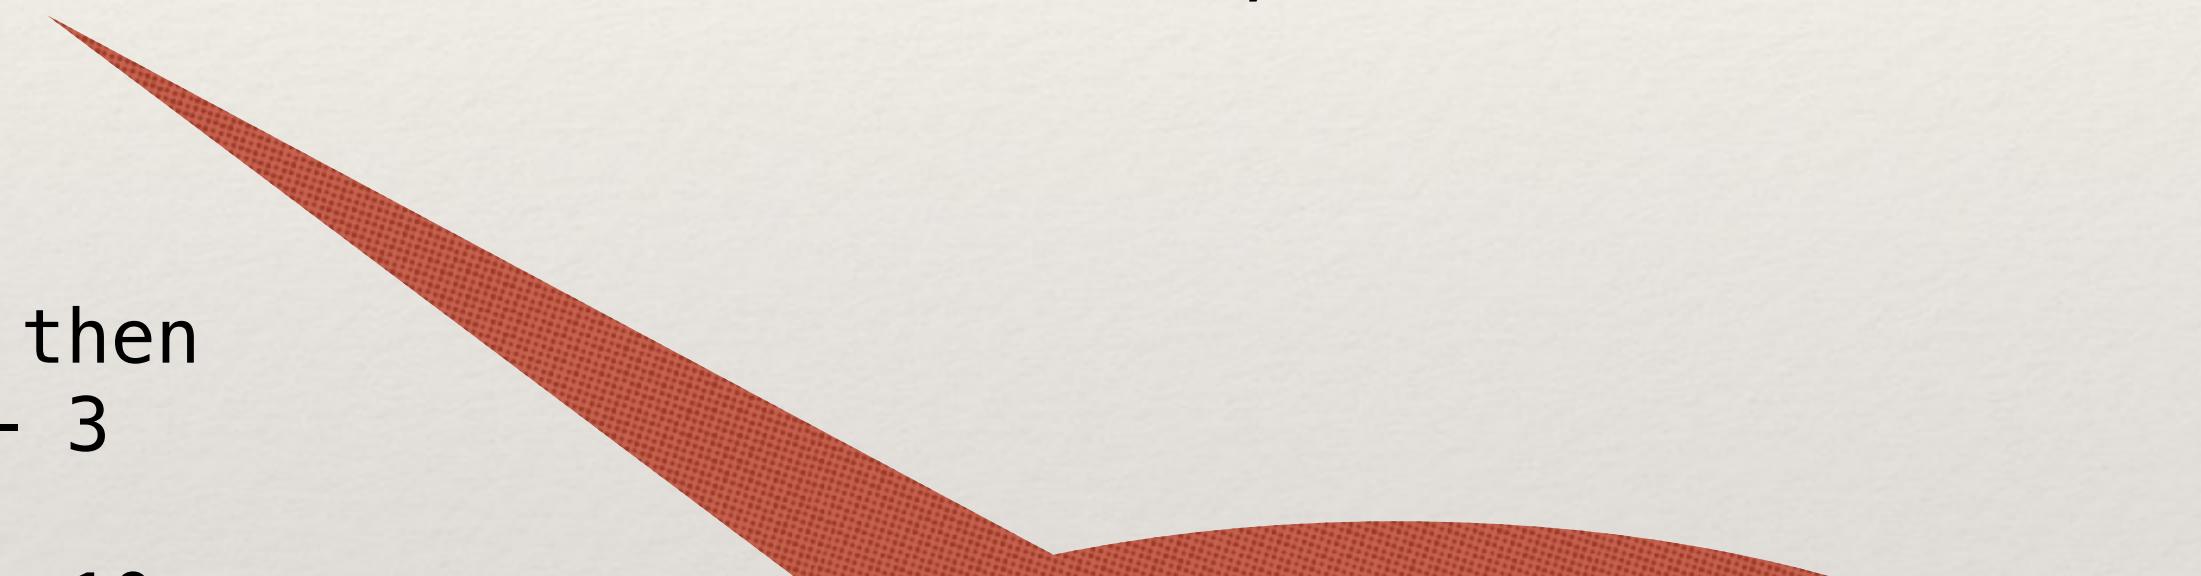
- ❖ The ϕ instruction was designed to choose at that time which branch we just came from
- ❖ ϕ was so-named by IBM in the 80s when they developed the idea of SSA to stand for ‘phoney’, haha
- ❖ We insert ϕ to select the right name and assign it to a new name, which can then be used

```
let
  ni x1 is 42
in
  now x2 is x1 - 7;
  let
    ni y1 is 3
    ni w1 is 0
  in
    if x2 <= 35 then
      now y2 is x2 - 3
    else
      now y3 is x2 - 10
    end;
    now y4 is  $\phi$  (y2, y3)
    now w2 is y4 * 2;
    w2
  end
end
```

Translating to LLVM IR

```
let
  ni x1 is 42
in
  now x2 is x1 - 7;
  let
    ni y1 is 3
    ni w1 is 0
  in
    if x2 <= 35 then
      y2 is x2 - 3
    else
      y3 is x2 - 10
    end;
    now y4 is  $\phi$  (y2, y3)
    now w2 is y4 * 2;
    w2
  end
end
```

```
%x1 = add i64 42, 0
%x2 = sub i64 %x1, 7
%y1 = add i64 3, 0
%w1 = add i64 3, 0
```



We begin by the simple translation of the variables

Translating to LLVM IR

```
let
  ni x1 is 42
in
  now x2 is x1 - 7;
  let
    ni y1 is 3
    ni w1 is 0
  in
    if x2 <= 35 then
      y2 is x2 - 3
    else
      y3 is x2 - 10
    end;
    now y4 is  $\Phi$ (y2, y3)
    now w2 is y4 * 2;
    w2
  end
end
```

```
%x1 = add i64 42, 0
%x2 = sub i64 %x1, 7
%y1 = add i64 3, 0
%w1 = add i64 3, 0
```

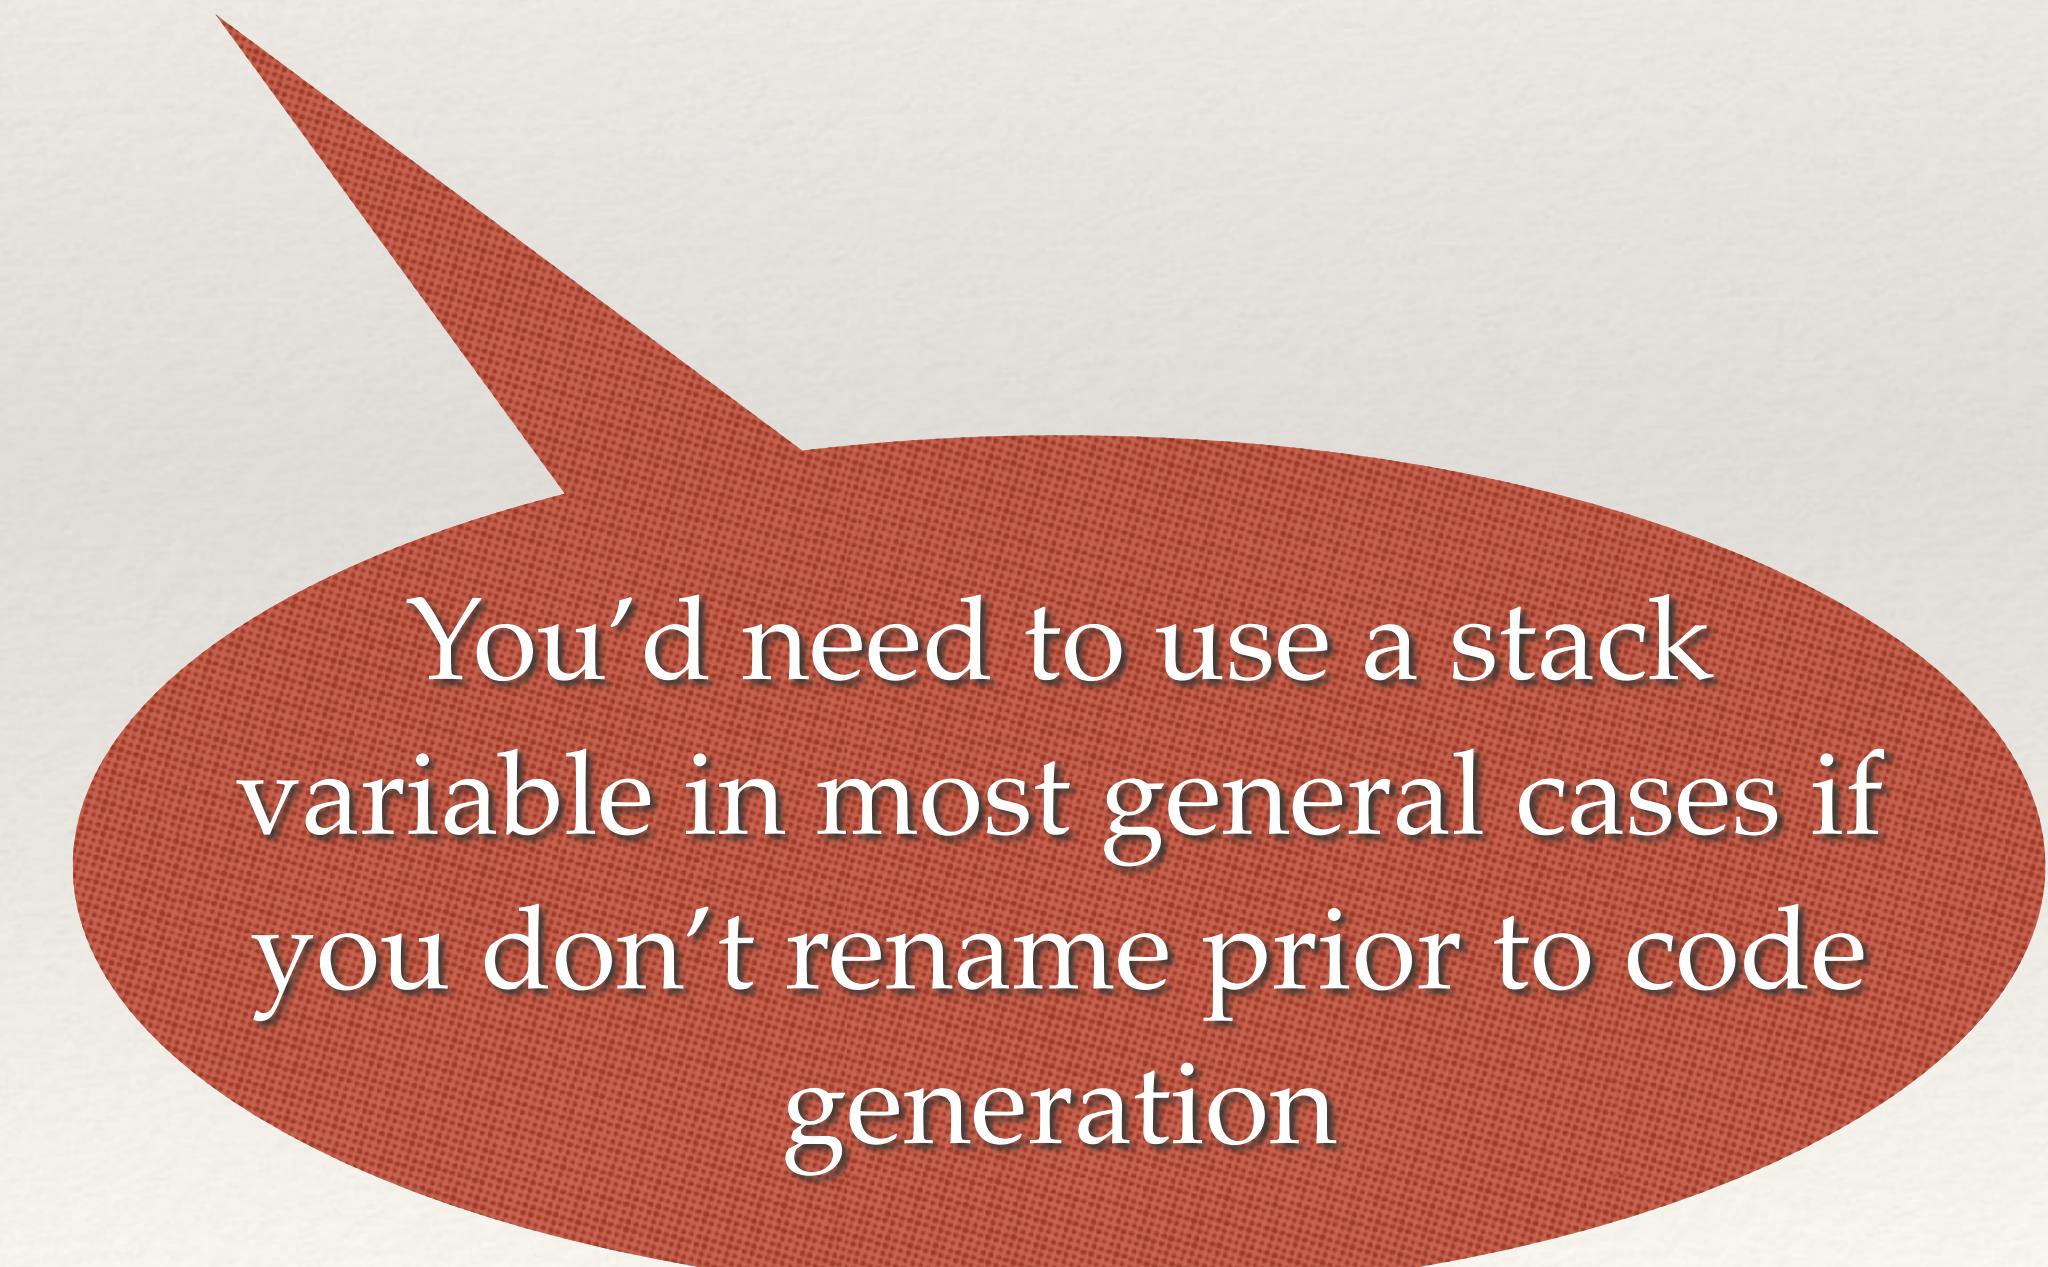


Note, we're just using register values, which assumes you renamed everything already

Translating to LLVM IR

```
let
  ni x1 is 42
in
  now x2 is x1 - 7;
  let
    ni y1 is 3
    ni w1 is 0
  in
    if x2 <= 35 then
      y2 is x2 - 3
    else
      y3 is x2 - 10
    end;
    now y4 is  $\phi$  (y2, y3)
    now w2 is y4 * 2;
    w2
  end
end
```

```
%x1 = add i64 42, 0
%x2 = sub i64 %x1, 7
%y1 = add i64 3, 0
%w1 = add i64 3, 0
```



You'd need to use a stack variable in most general cases if you don't rename prior to code generation

Translating a Conditional Branch

```
let
  ni x1 is 42
in
  now x2 is x1 - 7;
  let
    ni y1 is 3
    ni w1 is 0
  in
    if x2 <= 35 then
      y2 is x2 - 3
    else
      y3 is x2 -
    end;
  now y4 is (y2, y3)
    4 * 2;
```

```
%x1 = add i64 42, 0
%x2 = sub i64 %x1, 7
%y1 = add i64 3, 0
%w1 = add i64 3, 0
; if-then test
%t1 = icmp sle i64 %x2, 35
br i1 %t1, label %L1, label %L2
L1:
  %y2 = sub i64 %x2, 3
  br label L3
L2:
  %y3 = sub i64 %x2, 10
  br label L3
L3:
  %y4 = phi i64 [ %y2, L1 ], [ %y3, L2 ]
  %w2 = mul i64 %y4, 2
```

Next we encode
the loop using LLVM's
branch

Translating a Conditional Branch

```
let
  ni x1 is 42
in
  now x2 is x1 - 7;
  let
    ni y1 is 3
    ni w1 is 0
  in
    if x2 <= 35 then
      y2 is x2 - 3
    else
      y3 is x2 - 10
    end;
  now y4 is (y2 + y3) * 2;
```

This requires using icmp first, followed by a br i1 test

```
%x1 = add i64 42, 0
%x2 = sub i64 %x1, 7
%y1 = add i64 3, 0
%w1 = add i64 3, 0
; if-then test
%t1 = icmp sle i64 %x2, 35
br i1 %t1, label %L1, label %L2
L1:
  %y2 = sub i64 %x2, 3
  br label L3
L2:
  %y3 = sub i64 %x2, 10
  br label L3
L3:
  %y4 = phi i64 [ %y2, L1 ], [ %y3, L2 ]
  %w2 = mul i64 %y4, 2
```

Translating a Conditional Branch

```
let
  ni x1 is 42
in
  now x2 is x1 - 7;
  let
    ni y1 is 3
    ni w1 is 0
  in
    if x2 <= 35 then
      y1 is x2 - 3
    else
      y1 is 10
    y3)
;

%x1 = add i64 42, 0
%x2 = sub i64 %x1, 7
%y1 = add i64 3, 0
%w1 = add i64 3, 0
; if-then test
%t1 = icmp sle i64 %x2, 35
br i1 %t1, label %L1, label %L2
L1:
  %y2 = sub i64 %x2, 3
  br label L3
L2:
  %y3 = sub i64 %x2, 10
  br label L3
L3:
  %y4 = phi i64 [ %y2, L1 ], [ %y3, L2 ]
  %w2 = mul i64 %y4, 2
```

Notice that for the conditional branch, we must specify labels for where we branch to on true and false

Translating an Unconditional Branch

```
let
  ni x1 is 42
in
  now x2 is x1 - 7;
  let
    ni y1 is 3
    ni w1 is 0
  in
    if x2 <= 35 then
      v2 is x2 - 3
    else
      v2 is x2 + 10
    in
      y2 is v2 * 3
      y3 is v2 * 10
    ;
```

Furthermore,
notice that we have an
unconditional branch to exit
our different branches.

These start with “br
label”

```
%x1 = add i64 42, 0
%x2 = sub i64 %x1, 7
%y1 = add i64 3, 0
%w1 = add i64 3, 0
; if-then test
%t1 = icmp sle i64 %x2, 35
br i1 %t1, label %L1, label %L2
L1:
  %y2 = sub i64 %x2, 3
  br label L3
L2:
  %y3 = sub i64 %x2, 10
  br label L3
L3:
  %y4 = phi i64 [ %y2, L1 ], [ %y3, L2 ]
  %w2 = mul i64 %y4, 2
```

Using ϕ

Finally, we
translate the ϕ
instruction

```
if x2 is 35 then
    y2 is x2 - 3
else
    y3 is x2 - 10
end;
now y4 is  $\phi$  (y2, y3)
now w2 is y4 * 2;
w2
end
end
```

```
%x1 = add i64 42, 0
%x2 = sub i64 %x1, 7
%y1 = add i64 3, 0
%w1 = add i64 3, 0
; if test
%t1 = icmp sle i64 %x2, 35
br i1 %t1, label %L1, label %L2
L1:
    %y2 = sub i64 %x2, 3
    br label L3
L2:
    %y3 = sub i64 %x2, 10
    br label L3
L3:
    %y4 = phi i64 [ %y2, L1 ], [ %y3, L2 ]
    %w2 = mul i64 %y4, 2
```

Using Φ

```
let
  ni x1 is 42
in
  now x2 is x1 - 7;
  let
    ni y1 is 3
    ni w1 is 0
  in
    if x2 <= 35 then
      y2 is x2 - 3
    else
      y3 is x2 - 10
    end;
    now y4 is  $\Phi$  (y2, y3)
    now w2 is y4 * 2;
    w2
  end
end
```

```
%x1 = add i64 42, 0
%x2 = sub i64 %x1, 7
%y1 = add i64 3, 0
%w1 = add i64 3, 0
; if test
%t1 = icmp sle i64
br i1 %t1, label %L1
L1:
  %y2 = sub i64 %x2, 3
  br label L3
L2:
  %y3 = sub i64 %x2, 10
  br label L3
L3:
  %y4 = phi i64 [ %y2, L1 ], [ %y3, L2 ]
  %w2 = mul i64 %y4, 2
```

This becomes
“phi”, and has a list of
first class typed values
and Labels

Using ϕ

```
let
  ni x1 is 42
in
  now x2 is x1 - 7;
  let
    ni y1 is 3
    ni w1 is 0
  in
    if x2 <= 35 then
      y2 is x2 - 3
    else
      y3 is x2 - 10
    end;
    now y4 is  $\phi$  (y2, y3)
    now w2 is y4 * 2;
    w2
  end
end
```

```
%x1 = add i64 42
%x2 = sub i64
%y1 = add i64
%w1 = add i64
; if test
%t1 = icmp
br i1 %t1,
L1:
  %y2 = sub i64
  br label L3
L2:
  %y3 = sub i64 %x2, 10
  br label L3
L3:
  %y4 = phi i64 [ %y2, L1 ], [ %y3, L2 ]
  %w2 = mul i64 %y4, 2
```

This list can be longer than two tests, and each pair is the value to copy for assignment, and the label of the predecessor block

Using ϕ

```
let
  ni x1 is 42
in
  now x2 is x1 - 7;
  let
    ni y1 is 3
    ni w1 is 0
  in
    if x2 <= 35 then
      y2 is x2 - 3
    else
      y3 is x2 - 10
    end;
    now y4 is  $\Phi$  (y2, y3)
    now w2 is y4 * 2;
    w2
  end
end
```

```
%x1 = add i64 42, 0
%x2 = sub i64 %x1, 7
%y1 = add i64 3, 0
%w1 = add i64 3, 0
; if test
%t1 = icmp sle i64 %x2, 35
br i1 %t1, label %L1, label %L2
L1:
  %y2 = sub i64 %x2, 3
  br label L3
L2:
  %y3 = sub i64 %x2, 10
  br label L3
L3:
  %y4 = phi i64 [ %y2, L1 ], [ %y3, L2 ]
  %w2 = mul i64 %y4, 2
```

And now we can
simply use y4 at this
point forward

Translating Binary Operators

- ❖ We've already seen how to use some binary operators
- ❖ The integer arithmetic ones are as expected: add, sub, mul, sdiv
- ❖ These take a type, and two operands which must be first class types

```
%x1 = add i64 42, 0
%x2 = sub i64 %x1, 7
%y1 = add i64 3, 0
%w1 = add i64 3, 0
; if test
%t1 = icmp sle i64 %x2, 35
br i1 %t1, label %L1, label %L2
L1:
%y2 = sub i64 %x2, 3
br label L3
L2:
%y3 = sub i64 %x2, 10
br label L3
L3:
%y4 = phi i64 [ %y2, L1 ], [ %y3, L2 ]
%w2 = mul i64 %y4, 2
```

Translating Binary Operators

- ❖ Note that as a computation, they must be assigned to a local of some sort
- ❖ This can then be stored in a stack value if needed

```
%x1 = add i64 42, 0
%x2 = sub i64 %x1, 7
%y1 = add i64 3, 0
%w1 = add i64 3, 0
; if test
%t1 = icmp sle i64 %x2, 35
br i1 %t1, label %L1, label %L2
L1:
%y2 = sub i64 %x2, 3
br label L3
L2:
%y3 = sub i64 %x2, 10
br label L3
L3:
%y4 = phi i64 [ %y2, L1 ], [ %y3, L2 ]
%w2 = mul i64 %y4, 2
```

Translating Binary Operators

- ❖ We also have bitwise binary operators:
 - ❖ shl (for logical shift left)
 - ❖ lshr (for logical shift right)
 - ❖ ash (for arithmetic shift right, which keeps the signed value)
 - ❖ and (logical and)
 - ❖ or (logical or)
 - ❖ xor (logical xor)

```
%x1 = add i64 42, 0
%x2 = sub i64 %x1, 7
%y1 = add i64 3, 0
%w1 = add i64 3, 0
; if test
%t1 = icmp sle i64 %x2, 35
br i1 %t1, label %L1, label %L2
L1:
    %y2 = sub i64 %x2, 3
    br label L3
L2:
    %y3 = sub i64 %x2, 10
    br label L3
L3:
    %y4 = phi i64 [ %y2, L1 ], [ %y3, L2 ]
    %w2 = mul i64 %y4, 2
```

Translating Functions

- ❖ Functions are fairly easy to translate to LLVM IR because they follow a similar definition as to how they are displayed in Ni.
- ❖ Here we have a simple Ni language function

```
// add one to x  
neewom incr(int x) as int is  
  x + 1
```

Translating Functions

- ❖ To create a function, we use *define*

```
// add one to x
neewom incr(int x) as int is
    x + 1

// add one to x
define i64 @incr(i64 %x) nounwind {
entry:
    %1 = alloca i64, align 8
    store i64 %x, %1 align 8
    %2 = load i64* %1, align 8
    %3 = add nsw i64 %2, 1
    ret i64 %3
}
```

Translating Functions

- ❖ To create a function, we use *define*
- ❖ For Ni, all functions will also have the *nounwind* attribute—this means they do not throw exceptions

```
// add one to x
neewom incr(int x) as int is
    x + 1

// add one to x
define i64 @incr(i64 %x) nounwind {
entry:
    %1 = alloca i64, align 8
    store i64 %x, %1 align 8
    %2 = load i64* %1, align 8
    %3 = add nsw i64 %2, 1
    ret i64 %3
}
```

Translating Functions

- ❖ As with Ni, we declare the return type

```
// add one to x
neewom incr(int x) as int is
    x + 1

// add one to x
define i64 @incr(i64 %x) nounwind {
entry:
    %1 = alloca i64, align 8
    store i64 %x, %1 align 8
    %2 = load i64* %1, align 8
    %3 = add nsw i64 %2, 1
    ret i64 %3
}
```

Translating Functions

- ❖ As with Ni, we declare the return type
- ❖ And the name of the function is a global, while the formal parameter is a local

```
// add one to x
neewom incr(int x) as int is
    x + 1

// add one to x
define i64 @incr(i64 %x) nounwind {
entry:
    %1 = alloca i64, align 8
    store i64 %x, %1 align 8
    %2 = load i64* %1, align 8
    %3 = add nsw i64 %2, 1
    ret i64 %3
}
```

Translating Functions

- ❖ Every function in LLVM has an entry block, whether you label it or not

```
// add one to x
neewom incr(int x) as int is
    x + 1

// add one to x
define i64 @incr(i64 %x) nounwind {
entry:
    %1 = alloca i64, align 8
    store i64 %x, %1 align 8
    %2 = load i64* %1, align 8
    %3 = add nsw i64 %2, 1
    ret i64 %3
}
```

Translating Functions

- ❖ Every function in LLVM has an entry block, whether you label it or not
- ❖ Unlike other labels, *you cannot jump to the first block in a function*

```
// add one to x
neewom incr(int x) as int is
    x + 1

// add one to x
define i64 @incr(i64 %x) nounwind {
entry:
    %1 = alloca i64, align 8
    store i64 %x, %1 align 8
    %2 = load i64* %1, align 8
    %3 = add nsw i64 %2, 1
    ret i64 %3
}
```

Translating Functions

- ❖ This is simply because on entry to a function, LLVM must emit certain code—if you need to jump to the entry, what you really need is a *call*, i.e., a function call

```
// add one to x
neewom incr(int x) as int is
    x + 1

// add one to x
define i64 @incr(i64 %x) nounwind {
entry:
    %1 = alloca i64, align 8
    store i64 %x, %1 align 8
    %2 = load i64* %1, align 8
    %3 = add nsw i64 %2, 1
    ret i64 %3
}
```

Translating Functions

- ❖ Note that the first thing we do is copy the formal parameters into temporaries—this is because we have no idea if the function will be recursive
- ❖ Also note that formal parameters look like they're on the stack by default

```
// add one to x
neewom incr(int x) as int is
    x + 1

// add one to x
define i64 @incr(i64 %x) nounwind {
entry:
    %1 = alloca i64, align 8
    store i64 %x, %1 align 8
    %2 = load i64* %1, align 8
    %3 = add nsw i64 %2, 1
    ret i64 %3
}
```

Translating Functions

- ❖ We allocate space for `%x`

```
// add one to x
neewom incr(int x) as int is
    x + 1

// add one to x
define i64 @incr(i64 %x) nounwind {
entry:
    %1 = alloca i64, align 8
    store i64 %x, %1 align 8
    %2 = load i64* %1, align 8
    %3 = add nsw i64 %2, 1
    ret i64 %3
}
```

Translating Functions

- ❖ We allocate space for `%x`
- ❖ Then we copy it to our local stack space

```
// add one to x
neewom incr(int x) as int is
    x + 1

// add one to x
define i64 @incr(i64 %x) nounwind {
entry:
    %1 = alloca i64, align 8
    store i64 %x, %1 align 8
    %2 = load i64* %1, align 8
    %3 = add nsw i64 %2, 1
    ret i64 %3
}
```

Translating Functions

- ❖ We allocate space for `%x`
- ❖ Then we copy it to our local stack space
- ❖ Then we move it to a register so we can operate on it

```
// add one to x
neewom incr(int x) as int is
    x + 1

// add one to x
define i64 @incr(i64 %x) nounwind {
entry:
    %1 = alloca i64, align 8
    store i64 %x, %1 align 8
    %2 = load i64* %1, align 8
    %3 = add nsw i64 %2, 1
    ret i64 %3
}
```

Translating Functions

- ❖ We allocate space for `%x`
- ❖ Then we copy it to our local stack space
- ❖ Then we move it to a register so we can operate on it
- ❖ We operate on it, store the result

```
// add one to x
neewom incr(int x) as int is
    x + 1

// add one to x
define i64 @incr(i64 %x) nounwind {
entry:
    %1 = alloca i64, align 8
    store i64 %x, %1 align 8
    %2 = load i64* %1, align 8
    %3 = add nsw i64 %2, 1
    ret i64 %3
}
```

Translating Functions

- ❖ And return it using a *ret* instruction

```
// add one to x
neewom incr(int x) as int is
    x + 1

// add one to x
define i64 @incr(i64 %x) nounwind {
entry:
    %1 = alloca i64, align 8
    store i64 %x, %1 align 8
    %2 = load i64* %1, align 8
    %3 = add nsw i64 %2, 1
    ret i64 %3
}
```

Recursive Functions

```
// add one to x
newom sum(int n) as int is
    if n == 0 then
        0
    else
        1 + sum(n - 1)
end
```

- ❖ Recursive functions differ little from regular functions, except that you have a function call

Recursive Functions

```
// add one to x
newom sum(int n) as int is
    if n == 0 then
        0
    else
        1 + sum(n - 1)
end
```

```
define i64 @sum(i64 %n) nounwind {
entry:
%0 = icmp eq i64 %n, 0
br i1 %0, label %exit, label %recurse

recurse:
%1 = sub i64 %n, 1
%2 = call i64 @sum(i64 %1)
%3 = add i64 %2, %n
ret i64 %3

exit:
ret i64 0
}
```

We define a label for the entry into our function

- ❖ Recursion is just like a function call

from regular functions, except that you have

Recursive Functions

```
// add one to x
newom sum(int n) as int is
    if n == 0 then
        0
    else
        1 + sum(n - 1)
end
```

```
define i64 @sum(i64 %n) nounwind {
entry:
    %0 = icmp eq i64 %n, 0
    br i1 %0, label %exit, label %recurse

recurse:
    %1 = sub i64 %n, 1
    %2 = call i64 @sum(i64 %1)
    %3 = add i64 %2, %n
    ret i64 %3

exit:
    ret i64 0
}
```

We start with the branch, if n is 0, we conditionally branch to a function.

- ❖ Recursive functions work just like regular functions, except that you have a function that calls itself.

Recursive Functions

```
// add one to x
newom sum(int n) as int is
    if n == 0 then
        0
    else
        1 + sum(n - 1)
end
```

```
define i64 @sum(i64 %n) nounwind {
entry:
    %0 = icmp eq i64 %n, 0
    br i1 %0, label %exit, label %recurse
recurse:
    %1 = sub i64 %n, 1
    %2 = call i64 @sum(i64 %1)
    %3 = add i64 %2, %n
    ret i64 %3
exit:
    ret i64 0
}
```

We have two options to branch, to the exit (because n is 0) or to the recursive part, which we label

- ❖ Recursion is just like regular functions, except that you have a function that calls itself

Recursive Functions

```
// add one to x
newom sum(int n) as int is
    if n == 0 then
        0
    else
        1 + sum(n - 1)
end
```

```
define i64 @sum(i64 %n) nounwind {
entry:
    %0 = icmp eq i64 %n, 0
    br i1 %0, label %exit, label %recurse

recurse:
    %1 = sub i64 %n, 1
    %2 = call i64 @sum(i64 %1)
    %3 = add i64 %2, %n
    ret i64 %3

exit:
    ret i64 0
}
```

If n is not 0, we subtract 1

- ❖ Recurse from n, and call sum(n - 1). We return regular functions, except that you have a function.

Recursive Functions

```
// add one to x
newom sum(int n) as int is
    if n == 0 then
        0
    else
        1 + sum(n - 1)
end
```

```
define i64 @sum(i64 %n) nounwind {
entry:
    %0 = icmp eq i64 %n, 0
    br i1 %0, label %exit, label %recurse

recurse:
    %1 = sub i64 %n, 1
    %2 = call i64 @sum(i64 %1)
    %3 = add i64 %2, %n
    ret i64 %3

exit:
    ret i64 0
}
```

- ❖ Recursion is similar to regular functions, except that you have a function that calls itself.

Recursive Functions

```
// add one to x
newom sum(int n) as int is
    if n == 0 then
        0
    else
        1 + sum(n - 1)
end
```

```
define i64 @sum(i64 %n) nounwind {
entry:
    %0 = icmp eq i64 %n, 0
    br i1 %0, label %exit, label %recurse

recurse:
    %1 = sub i64 %n, 1
    %2 = call i64 @sum(i64 %1)
    %3 = add i64 %2, %n
    ret i64 %3

exit:
    ret i64 0
}
```

Note that we didn't put this into tail-recursive form, which we could have

- ❖ Recursive functions work similarly to regular functions, except that you have a function that calls itself.

Recursive Functions

```
// add one to x
newom sum(int n) as int is
    if n == 0 then
        0
    else
        1 + sum(n - 1)
end
```

```
define i64 @sum(i64 %n) nounwind {
entry:
    %0 = icmp eq i64 %n, 0
    br i1 %0, label %exit, label %recurse

recurse:
    %1 = sub i64 %n, 1
    %2 = call i64 @sum(i64 %1)
    %3 = add i64 %2, %n
    ret i64 %3

exit:
    ret i64 0
}
```

- ❖ Recursive functions differ little from regular functions, except that you have a function call

References

- ❖ <http://llvm.org/docs/LangRef.html>
- ❖ <http://www.aosabook.org/en/llvm.html>
- ❖ <http://llvm.lyngvig.org/Articles/Mapping-High-Level-Constructs-to-LLVM-IR>